

A Unified Physics-Aware Quantum Machine Learning Framework across Power GaN HEMTs and Logic Nanowire FETs: Predicting Unseen Process Splits and Held-Out Geometry Combinations with Lower Error and Tighter Split-to-Split Variability

Rushat Rai¹, Yun-Yuan Wang², Autsada Kakaen³, Pei-Jie Chang¹, Doan Viet Nguyen¹, Yuan-Chieh Chiu¹, Doldet Tantraviwat³, Niall Tumilty¹, Simon See⁴, Wen-Jay Lee^{5,*}, Tai-Yue Li^{5,*}, Nan-Yow Chen^{5,*}, and Tian-Li Wu^{1,*}

¹ National Yang Ming Chiao Tung University, Taiwan. ² NVIDIA AI Technology Center, Taiwan. ³ Chiang Mai University, Thailand.

⁴ NVIDIA AI Technology Center, Singapore. ⁵ National Center for High-performance Computing, Taiwan.

*E-mail: {wjlee, 2503001, nanyow}@nlar.org.tw, tlwu@nycu.edu.tw

Abstract- We present a unified reinforcement-learning (RL) framework that discovers compact parametrized quantum circuits (PQCs) for data-scarce device modeling. A graph neural network (GNN) policy optimized by proximal policy optimization (PPO) searches circuit architectures using leave-one-group-out cross-validation (LOGOCV) error on held-out process or geometry groups as the reward. The framework achieves the lowest mean absolute error (MAE) on all 11 targets versus six classical baselines, with 59% lower error (I_{off}) and 81% tighter fold variability (V_{TH}) for HEMTs and 84% lower error (V_{TH} , SS, I_{off}) and 82% tighter fold variability (I_{off}) for NWFETs. These results demonstrate the potential of RL-selected, classically simulated PQCs as compact surrogates with low OOD error and improved physical consistency, despite imposing no explicit physical constraints, penalty terms, or device-specific equations, on the two evaluated device datasets.

I. INTRODUCTION

Advanced device architectures are evolving to meet the distinct demands of high-voltage power conversion and sub-3-nm logic. p-GaN gate GaN HEMTs enable normally-off operation with controllable threshold voltage for high-performance power switching [1], while gate-all-around nanowire FETs (NWFETs) suppress short-channel effects and off-state leakage for continued scaling toward low-power logic [2].

Their optimization is limited by costly fabrication and TCAD cycles and by stochastic variation from charge trapping, thermal gradients, and parameter extraction. High-value design exploration thus requires physically consistent extrapolation to unfabricated process splits and held-out geometry combinations, not merely interpolation among measured devices (**Fig. 1**).

Existing surrogates often embed device-specific equations in the architecture [16]–[18], impeding reuse across device families, or report only within-distribution tests [16][17]. We instead impose no explicit physical constraints, penalty terms, or device-specific equations, making out-of-distribution (OOD) prediction the explicit objective; physical consistency emerges from the OOD reward. Parametrized quantum circuits (PQCs) provide a compact Fourier-structured function class whose expressivity grows with depth and data re-uploading [7][8] rather than parameter count [9]. A relational graph-attention policy represents candidate circuits as heterogeneous graphs, optimized by PPO with a leave-one-group-out cross-validation reward, selecting architectures for unseen process or geometry groups rather than average training fit.

A single search framework applies to both planar HEMTs and 3D NWFETs without re-deriving device equations, discovering a circuit optimized for each (**Fig. 4**): HEMTs use physics-based pretraining followed by calibration to measured wafer data, while NWFETs learn directly from 3D TCAD. Across 11 targets, the selected circuits reduce OOD error by 6–34% for HEMTs and 25–39% for NWFETs, with about 5× tighter HEMT split-to-split variability. To our knowledge, this is the first experimentally validated, RL architecture-searched PQC framework [10] for

multi-target device modeling optimized directly for OOD prediction (**Table 3**).

II. METHODOLOGY

Dataset preparation: To simulate p-GaN HEMTs for pretraining, a 1D electrostatic solver incorporating Mg-acceptor ionization, polarization charge, 2D electron-gas (2DEG) formation, and interface traps [3]–[5] (**Fig. 5**) was developed. Solved by finite-volume discretization with damped Newton iteration, it extracts full transfer curves and six targets: forward/reverse threshold voltages ($V_{\text{TH,for}}/V_{\text{TH,rev}}$), hysteresis (ΔV_{TH}), subthreshold swing (ss), on-current (I_{ON}), and off-current (I_{OFF}). Electrostatic analysis confirms physical validity (**Fig. 3**): smooth band bending, gate-modulated surface potential, enhancement-mode operation with expected 2DEG densities, and charge conservation at the sweep turnaround. The solver generates 1,785 devices/h on a CPU (~90× faster than TCAD), while a PhysicsNeMo PINN reaches 12,898 devices/h on a GPU (~650× faster) [6] (**Table 3**) to augment it. The experimental HEMT dataset comprises 544 devices spanning unique process splits and wafer positions; the NWFET dataset comprises 1,914 3D-TCAD devices for five targets across eight groups defined by channel length, nanowire radius, and dielectric thickness.

Encoding and circuit search: Device features map one-to-one to qubits by angle encoding: five process/position inputs for HEMTs and three geometry inputs (channel length (L_g), radius (r), dielectric thickness (t)) for NWFETs. The agent selects a pre-encoder, R_x , R_y , or R_z rotations, and entanglers from CNOT, CRX, CRY, and CRZ families arranged as chain, ring, all-to-all, or even/odd pairs. The shared library contains 5 rotation and 24 entangling blocks plus data re-uploading, with candidate depth 3–12 blocks and at most four re-uploads. Validity masks require at least one rotation and one entangler, forbid consecutive re-uploads, and allow early termination, limiting unproductive exploration.

Readout: Single-qubit Z expectations pass through a trainable tempered tanh and linear map to standardized multi-target outputs. For measured HEMTs, a wafer-position correction is added to the quantum prediction, letting the PQC capture process-dependent behavior while the small residual head accounts for spatial drift across the wafer.

GNN policy and RL: Each partial circuit is encoded as a typed graph with feature, qubit, operation-block, and readout nodes linked by circuit-specific relations (**Fig. 2**). A relational graph attention network serves as a shared actor-critic backbone [11], pooling masked node embeddings with global circuit features. The actor selects the pre-encoder, feature encoding, next valid block, or termination; the critic estimates terminal reward. PPO with generalized advantage estimation updates the policy [12][13], and a 0.001 per-block penalty discourages unnecessary depth. This preserves connectivity across variable-length candidates and device-specific qubit counts.

Training and selection. HEMT training uses three stages: synthetic electrostatic pretraining, wafer-position-head calibration with the PQC frozen, and joint experimental fine-tuning; NWFETs

use direct single-stage TCAD training. Architecture selection follows nested LOGOCV: the inner reward uses the worst validation-group error with complexity and invalid-structure penalties, so a strong average cannot hide failure on one split; the selected candidate is evaluated on an untouched outer group. Hybrid quantum-classical training minimizes mean-squared error alone, with no physics-informed loss or constraint terms, using parameter-shift gradients [14] on NVIDIA CUDA-Q [15].

III. RESULTS AND DISCUSSION

We compare our RL-selected hybrid quantum neural network (PQC) against six classical baselines on two OOD challenges: measured p-GaN HEMT process variation and 3D NWFET geometry scaling (Fig. 7). All models use the same group-wise validation protocol to prevent data leakage; results are averaged over 5 random seeds.

To isolate the model-class effect, the ANN baseline used the identical data pipeline, wafer-position correction, training stages, and LOGOCV splits as the PQC. It was swept across shallow-to-deep capacities plus a Fourier-feature ablation, and we report its best variant.

Benchmark accuracy and efficiency: The PQC has the lowest OOD MAE on all 11 targets versus the six classical baselines. For HEMTs, $V_{TH,for}$ MAE is 18.6% below the identically trained ANN and 48.0% below the worst baseline (PLS); $V_{TH,rev}$ MAE is 25.2% below the ANN and 42.1% below the worst baseline; and I_{OFF} MAE is 23.5% below the ANN and 59% below the worst baseline (Elastic Net). ΔV_{TH} , SS, and I_{ON} MAE also improve, while forward threshold-voltage fold variability is 53% tighter than the ANN and 81% tighter ($5.3\times$) than the worst baseline (Decision Tree). For NWFETs, the PQC leads on all five targets, V_{TH} , SS, and I_{OFF} MAE are 21.4%, 29.7%, and 37.5% below the ANN, respectively, and about 84% below the worst baseline (PLS); I_{OFF} fold variability is 82% tighter than the worst baseline (Fig. 13). It uses 2,460 trainable model parameters ($4.7\times$ fewer than the ANN) and sits on the best error-parameter Pareto frontier (Fig. 8).

Physical consistency in HEMTs: Although PQC and ANN have similar aggregate error distributions, the PQC better preserves threshold-hysteresis closure (Fig. 9), even though no hysteresis-closure constraint or physics-based penalty is applied during training, so this consistency arises from the OOD-rewarded architecture search rather than explicit enforcement. Consistency metrics pool out-of-fold predictions from models refitted per outer group. We measure the absolute deviation of ($V_{th,rev} - V_{th,forw}$) from the independently predicted hysteresis (ΔV_{TH}), which should be zero for consistent outputs. The PQC median consistency error is about $5.4\times$ lower than the ANN's; the ANN's inconsistent tail reaches roughly $4\times$ its own median (about $22\times$ the PQC median). Residual decomposition shows recipe-level mean bias falls to about 0.12 interquartile range (IQR) versus about 0.4 IQR for ANN, whereas die-to-die pattern error stays similar at about 0.4-0.5 IQR. The improvement mainly reflects suppressed wafer-to-wafer drift without losing local variation.

Retention of pretrained physics after fine-tuning: To test whether calibration preserves the process behavior acquired during solver pretraining, the fine-tuned process cores were reevaluated on the synthetic corpus with the wafer-position residual head excluded (Fig. 11). Before fine-tuning, both models retained all 16 evaluable solver-trend signs, with mean Spearman agreement $\rho_s = 0.949$ (ANN) and 0.935 (PQC). After fine-tuning, ANN sign retention fell to 75% and mean agreement to $\rho_s = 0.501$, reversing all four I_{on} trends, whereas PQC retained all 16 signs with mean $\rho_s = 0.875$. PQC also exceeded ANN in sign retention and trend

agreement across all LOGO folds, with substantially lower fine-tuning-induced degradation.

Joint-geometry OOD generalization in NWFETs: To evaluate electrostatic gate control across subthreshold and on-state regimes, we model three metrics: off-current, subthreshold swing, and maximum transconductance. For these radius-sensitive targets, the PQC reduces out-of-fold MAE relative to the ANN by 50.1%, 46.3%, and 39.2%, respectively. Because weak-inversion current is exponentially sensitive to small potential-barrier errors in held-out geometries, we bin normalized MAE by distance to the nearest training geometry (Fig. 10). Both models degrade deeper into OOD space, but the PQC degrades more slowly: its advantage grows from 45.5% (I_{OFF}) and 30.0% (SS) in the nearest bin to 51.6% and 58.9% in the deepest bin. The widening margin indicates the quantum prior captures geometry-dependent electrostatic scaling rather than memorizing sampled points. To verify that OOD gains reflect genuine physical consistency, we evaluate preservation of geometry-dependent TCAD behavior (Fig. 12). Along controlled slices with two geometry variables fixed, both models reproduce the dominant trends, but only the PQC tracks highly nonlinear responses such as the non-monotonic g_{max} dependence on dielectric thickness (t). Across all exact slices, the PQC achieves higher directional agreement with TCAD than the ANN (0.928 vs. 0.903) and a finite-difference slope ratio much closer to unity (0.979 vs. 0.877). The PQC also yields lower normalized MAE in 30 of 32 target-geometry combinations, with a median relative reduction of 42.2% and a worst-octant value of 0.133 vs. 0.333. Because these folds withhold joint low/high combinations of L_g , r , and t while using individual axis values seen elsewhere in training, the framework generalizes robustly to unseen joint geometries and missing design-space corners.

IV. CONCLUSION

We demonstrate the first experimentally validated, RL-searched PQC framework optimized directly for OOD device prediction. Table 4 highlights that, for the first time, the proposed framework uniquely combines multi-target OOD prediction, RL-searched quantum circuits, and scalability across distinct power and logic device families without embedding device-specific equations. A single physics-aware workflow generalizes across measured p-GaN HEMT process splits and unseen 3D-NWFET geometry combinations, achieving the lowest OOD MAE for all 11 targets among six classical baselines. Compared with the identically trained ANN, the PQC reduces MAE by up to 50.1% while using $4.7\times$ fewer parameters; versus all baselines, reductions reach 59% for HEMTs and 84% for NWFETs. Crucially, it preserves all 16 pretrained HEMT physical trends, improves threshold-hysteresis consistency by $5.4\times$, and maintains nonlinear NWFET geometry scaling without explicit physical constraints. These results establish OOD-rewarded quantum architecture search as a framework to discover parameter-efficient surrogate for data-scarce device optimization.

Acknowledgments: Supported by 1) NVIDIA Academic Grant Program and Dr. Chi-Cheng Fu (NVIDIA), 2) Advanced Semiconductor Technology Research Center under Higher Education Sprout Project of Ministry of Education, Taiwan, and 3) NSTC 114-2223-E-A49-003-MY4, 114-2119-M-007-013, 114-222-E-492-002-MY2. **References:** [1] C.-C. Tu et al., Nat. Rev. Electr. Eng., 2024. [2] Y.-J. Mii, IEDM, 2016. [3] B. Bakeroot et al., IEEE TED, 2018. [4] N. E. Posthuma et al., ISPSD, 2016. [5] L. Sayadi et al., IEEE TED, 2018. [6] M. Raissi et al., J. Comput. Phys., 2019. [7] M. Schuld et al., Phys. Rev. A, 2021. [8] A. Pérez-Salinas et al., Quantum, 2020. [9] M. C. Caro et al., Nat. Commun., 2022. [10] Y. Du et al., arXiv, 2020. [11] P. Veličković et al., ICLR, 2018. [12] J. Schulman et al., arXiv, 2017. [13] J. Schulman et al., ICLR, 2016. [14] M. Schuld et al., Phys. Rev. A, 2019. [15] J.-S. Kim et al., DAC, 2023. [16] Z. Zhang et al., IEDM, 2024. [17] Y.-M. Tseng et al., IEDM, 2025. [18] C. Zhang et al., IEDM, 2025. [19] A. Lu et al., EDTM, 2025. [20] Z. Wang et al., Adv. Sci., 2025.




Fig. 1. Emerging-device surrogate-modeling DTCO challenges and the proposed unified OOD modeling framework: (a) conventional DTCO workflow and modeling methods (Table 1); (b) high-value optimization requires physically consistent OOD extrapolation, where CML degrades; (c) target devices and data sources; and (d) data re-uploading POCs provide parameter-efficient expressivity scaling and enable a common framework that adapts across devices without re-deriving PDEs.

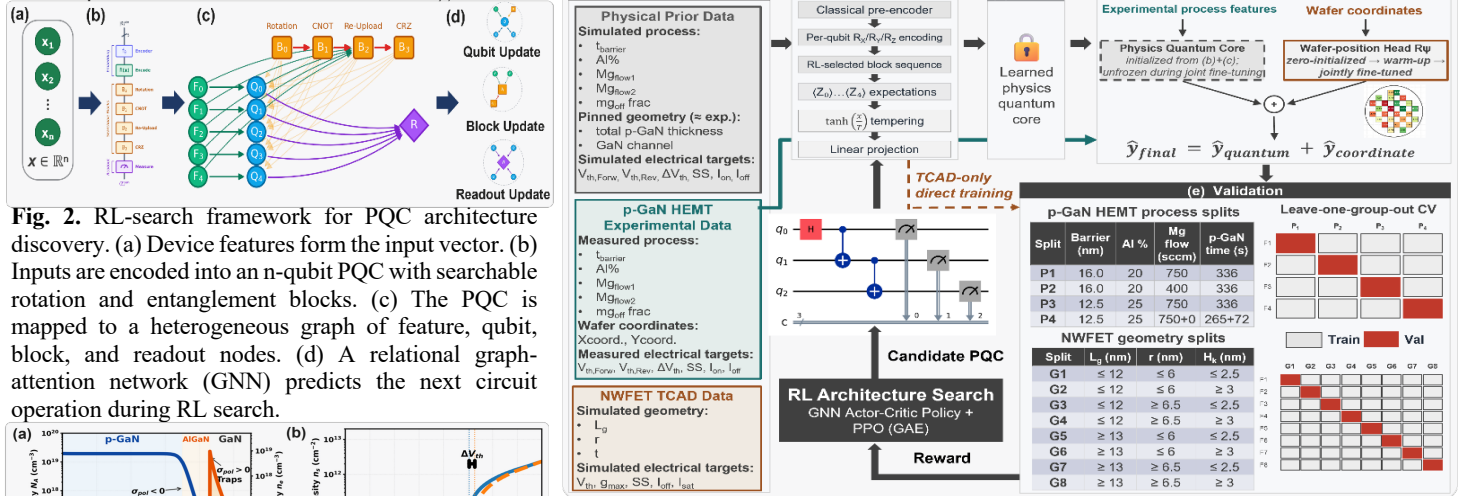


Fig. 2. RL-search framework for PQC architecture discovery. (a) Device features form the input vector. (b) Inputs are encoded into an n-qubit PQC with searchable rotation and entanglement blocks. (c) The PQC is mapped to a heterogeneous graph of feature, qubit, block, and readout nodes. (d) A relational graph-attention network (GNN) predicts the next circuit operation during RL search. (e) Validation results for p-GaN HEMT process splits and NWFET geometry splits using Leave-one-group-out CV.




Fig. 3. Electrostatic validation of the p-GaN HEMT solver: (a) doping profile, and (b) simulated 2DEG density versus gate bias.


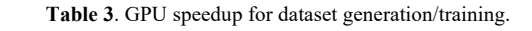










Fig. 5. Governing equations of the electrostatic solver for p-GaN HEMTs with considering the parameters in Table 2.

Table 2. Parameters used in the solver.

Parameter	Description	Value
ϵ_{GaIn}	Relative permittivity of GaIn	9.5
ϵ_{AlGaIn}	Relative permittivity of AlGaIn	9.5-0.5x
$P_{sp,GaN}$	Spontaneous polarization of GaN	-0.034 C/m ²
$P_{sp,AlN}$	Spontaneous polarization of AlN	-0.090 C/m ²
$P_{sp,AlGaIn}$	AlGaIn piezoelectric polarization model	-0.0525x+0.0282x ²
g_A	Mg acceptor degeneracy factor	4.0
N_V	Valence-band effective density of states	2.5×10 ¹⁹ cm ⁻³
N_C	Conduction-band effective density of states	2.2×10 ¹⁸ cm ⁻³
μ_0	Nominal low-field mobility prefactor	2000 cm ² /(V·s)
V_{DS}	Drain bias for low-field current	1.0 V
L_{gate}	Gate length	0.8 μm
$E_{g,GaN}$	GaN energy gap	3.4 eV
χ_{GaN}	GaN electron affinity	4.1 eV
$E_{g,AlN}$	AlN energy gap	6.2 eV
χ_{AlN}	AlN electron affinity	0.6 eV
χ_{bow}	AlGaIn bandgap bowing parameter	1.0 eV
N_{2D}	2-D density of states in GaN	8.36×10 ¹³ cm ⁻² eV ⁻¹
Φ_{Schott}	Schottky gate metal work function	4.8 eV
W_{trans}	Doping transition width	2.0 nm
λ_{tail}	Mg out-diffusion tail length	5.0 nm
α_{diff}	Mg out-diffusion tail fraction	0.02
λ_{env}	2DEG envelope decay length	3.0 nm
f_{act}	Fixed acceptor activation fraction	0.045
N_{st}	Slow-trap sheet density (sampled)	10 ¹¹ -10 ¹³ cm ⁻²
E_{offset}	Trap energy offset above E_C (sampled)	0.1-0.8 eV
τ_{tr}	Trap relaxation time constant (sampled)	10 ⁻² -10 ³ s
Mg coeffs	$N_A[\text{cm}^{-2}] = A[\text{Mg flow}[\text{sccm}]]^B$	A=3.251×10 ¹⁷ ; B=0.622



Approach	Hardware	Runtime	Throughput	Speedup
Pretraining		Dataset Generation		
Commercial TCAD simulation	CPU	days-months	~10-30 devices/h	1.0x
Physics-based solver	CPU	~8.4h	1,785 devices/h	~90x
PhysicsNeMo PINN	GPU	~1.2h	12,898 devices/h	~650x
Analytical Parameter-shift Model Training				
Standard Library	GPU	~8h	0.076 steps/s	1.0x
CUDA-Q	GPU	~0.3h	2.17 steps/s	~29x

**Table 4.** Comparison with related works

Ref	Dataset	Multi Target	OOD	Quantum component	Physics integration	Scalability
[16]	MTJ experimental + physical model	Yes	No	No	Physical equations embedded in GNN/DNN	Within MTJ/JMRAM family
[17]	FeFET TCAD + experimental	Yes	No	No	FeFET physics features embedded in GNN	Within FeFET family; UTBSOI with fine-tuning
[18]	SiC/GaN TCAD + SPICE simulations	Yes	Yes	No	Electro-thermal PDEs embedded in PINN	Across WBG families; SiC and GaN devices
[19]	Si NWFET TCAD	Yes	Yes	No	Data-trained physics-informed module	No
[20]	GaN HEMT experimental	No	No	Quantum kernel	No	No
Ours	p-GaN HEMT experimental + physical model; NWFET TCAD	Yes	Yes	RL-searched VQC	No device-specific equations embedded	Across distinct device families